

Intelligent Agents	State Representation Invariant	Informed Search Algorithms (2)	Simulated Annealing
<b>PEAS Framework</b> <ul style="list-style-type: none"> <li>Performance Measure</li> <li>Environment</li> <li>Actuators</li> <li>Sensors</li> </ul>	Abstract states (in the problem model) must have corresponding concrete states (in the real world).	<b>A* Search</b> <ul style="list-style-type: none"> <li>Evaluation function: Cost to reach <math>n</math> + Heuristic</li> <li>Time/space complexity: Same as DFS</li> <li>Complete?: Yes</li> <li>Optimal?: Depends on heuristic</li> <li>Heuristic = 0 <math>\rightarrow</math> UCS</li> </ul>	<ul style="list-style-type: none"> <li>Method to escape local minima</li> <li>Decrease temperature (randomness) slowly</li> </ul> <pre> current = initial state T = a large positive value while T &gt; 0:     next = a randomly selected successor of current     if value(next) &gt; value(current): current = next     else with probability P(current, next, T): current = next     decrease T return current </pre> <p>Allow "bad moves" from time to time</p> <p><math>P(\text{current}, \text{next}, T) = e^{\frac{\text{value}(\text{next}) - \text{value}(\text{current})}{T}}</math></p> <ul style="list-style-type: none"> <li>If T decreases slowly enough, high probability to find global optimum.</li> </ul>
Properties of Task Environment	Uninformed Search Algorithms	Admissible Heuristic	Adversarial Search
<ol style="list-style-type: none"> <li><b>Fully Observable</b> (vs. <b>Partially Observable</b>) <ul style="list-style-type: none"> <li>Agent has full environmental awareness</li> </ul> </li> <li><b>Deterministic</b> (vs. <b>Stochastic</b>) <ul style="list-style-type: none"> <li>Next environment state is only determined by agent's action</li> <li><b>Strategic</b>: If next environment state also depends on other agents</li> </ul> </li> <li><b>Episodic</b> (vs. <b>Sequential</b>) <ul style="list-style-type: none"> <li>Agent's actions depend solely on individual episodes.</li> </ul> </li> <li><b>Static</b> (vs. <b>Dynamic</b>) <ul style="list-style-type: none"> <li>Environment unchanged during agent deliberation; semi-dynamic if only performance score changes.</li> </ul> </li> <li><b>Discrete</b> (vs. <b>Continuous</b>) <ul style="list-style-type: none"> <li>Finite discrete actions</li> </ul> </li> <li><b>Single Agent</b> (vs. <b>Multi-agent</b>)</li> </ol>	<b>Breadth-first Search (BFS)</b> <i>b: branching factor, d: depth of solution</i> <ul style="list-style-type: none"> <li>Frontier: Queue</li> <li>Time Complexity: <math>O(b^d)</math></li> <li>Space Complexity: <math>O(b^d)</math></li> <li>Complete?: Yes, if <math>b</math> is finite</li> <li>Optimal?: Yes, if step cost is same everywhere</li> </ul> <b>Uniform-cost Search (UCS)</b> <i><math>\epsilon</math>: minimum edge cost, <math>C^*</math>: cost of optimal solution</i> <ul style="list-style-type: none"> <li>Frontier: Priority queue</li> <li>Time Complexity: <math>O(b^{\frac{C^*}{\epsilon}})</math></li> <li>Space Complexity: <math>O(b^{\frac{C^*}{\epsilon}})</math></li> <li>Complete?: Yes, if <math>\epsilon &gt; 0</math> and <math>C^*</math> finite</li> <li>Optimal?: Yes, if <math>\epsilon &gt; 0</math>  <math>\epsilon = 0</math> may cause zero cost cycle</li> </ul> <b>Depth-first Search (DFS)</b> <i>m: maximum depth</i> <ul style="list-style-type: none"> <li>Frontier: Stack</li> <li>Time Complexity: <math>O(b^m)</math></li> <li>Space Complexity: <math>O(bm)</math></li> <li>Complete?: No, when depth infinite or loops present</li> <li>Optimal?: No</li> </ul> <b>Depth-limited Search (DLS)</b> <i>l: depth limit of search</i> <ul style="list-style-type: none"> <li>Backtrack when limit is hit</li> <li>Time Complexity: <math>O(b^l)</math></li> <li>Space Complexity: <math>O(bl)</math></li> <li>Complete?: No</li> <li>Optimal?: No, if used with DFS</li> </ul> <b>Iterative Deepening Search (IDS)</b> <ul style="list-style-type: none"> <li>Time Complexity: <math>O(b^d)</math></li> <li>Space Complexity: <math>O(bd)</math>, if used with DFS</li> <li>Complete?: Yes</li> <li>Optimal?: Yes, if step cost is same everywhere</li> </ul>	$\forall n (h(n) \leq h^*(n))$ <p><math>h^*(n)</math>: true cost</p> <p>Admissible heuristic never over-estimates the cost to reach the goal. A* using <b>tree search</b> is optimal.</p>	Minimax
Structure of Agents	Informed Search Algorithms (1)	Consistent Heuristic	Properties: <ul style="list-style-type: none"> <li>Fully observable</li> <li>Deterministic</li> <li>Discrete</li> <li>Terminal states exist (no infinite runs)</li> </ul>
<b>Agent Function:</b> <ul style="list-style-type: none"> <li>completely specifies an agent.</li> <li>is <u>sufficient</u> to define an agent.</li> </ul> Common Agent Structures: <ul style="list-style-type: none"> <li>Simple reflex agents</li> <li>Model-based reflex agents</li> <li>Goal-based agents</li> <li>Utility-based agents</li> <li>Learning agents</li> </ul> <p style="text-align: right;">Increasing Complexity <math>\downarrow</math></p>	<b>Best-first Search</b> <ul style="list-style-type: none"> <li>Like UCS, but uses an evaluation function instead of step cost</li> </ul> <b>Greedy Best-first Search</b> <ul style="list-style-type: none"> <li>Evaluation function: Heuristic (cost of <math>n</math> to goal)</li> <li>Time/space complexity, completeness and optimality same as DFS</li> </ul>	$\forall n (h(n) \leq c(n, a, n') + h(n')) \wedge h(G) = 0$ <p><math>c(n, a, n')</math>: cost of <math>n</math> to <math>n'</math>  <math>n'</math>: a successor of <math>n</math></p> <p>Consistent heuristic is non-decreasing along any path. A* using <b>graph search</b> is optimal.</p>	<pre> def minimax(state):     v = max_value(state)     return action in successors(state) with value v  def max_value(state):     if is_terminal(state): return utility(state)     v = -∞     for action, next_state in successors(state):         v = max(v, min_value(next_state))     return v  def min_value(state):     if is_terminal(state): return utility(state)     v = ∞     for action, next_state in successors(state):         v = min(v, max_value(next_state))     return v </pre> <p>Assumes opponent play <b>optimally</b>: trying to <b>minimize</b> player's value</p> <ul style="list-style-type: none"> <li>Time Complexity: <math>O(b^m)</math></li> <li>Space Complexity: <math>O(bm)</math> with DFS</li> <li>Complete?: Yes, if tree finite</li> <li>Optimal?: Yes, against optimal opponent</li> </ul>
Exploration vs Exploitation	Hill Climbing		
<b>Exploration:</b> Discover new strategies or information that may lead to better long-term outcomes.	<ul style="list-style-type: none"> <li>Pick best among neighbours until no better neighbour</li> </ul>		
<b>Exploitation:</b> Maximize immediate rewards using known information.			
Problem Formulation			
<ul style="list-style-type: none"> <li>State representation</li> <li>Initial state</li> <li>Goal state / Goal test</li> <li>Actions</li> <li>Transition model</li> <li>Action cost function</li> </ul>			

## Alpha-beta Pruning

```
def alpha_beta_search(state):
    v = max_value(state, -∞, ∞)
    return action in successors(state) with value v

def max_value(state, α, β):
    if is_terminal(state): return utility(state)
    v = -∞
    for action, next_state in successors(state):
        v = max(v, min_value(next_state, α, β))
        α = max(α, v)
        if v >= β: return v
    return v

def min_value(state, α, β):
    if is_terminal(state): return utility(state)
    v = ∞
    for action, next_state in successors(state):
        v = min(v, max_value(next_state, α, β))
        β = min(β, v)
        if v <= α: return v
    return v
```

## Supervised Learning

**Regression:** Predict continuous output  
**Classification:** Predict discrete output

## Performance Measure

**For Regression:**

- Mean Squared Error

$$MSE = \frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2$$

- Mean Absolute Error

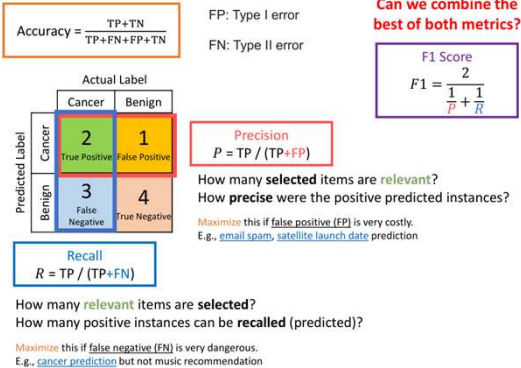
$$MAE = \frac{1}{N} \sum_{i=1}^N |\hat{y}_i - y_i|$$

**For Classification:**

- Correctness
- Accuracy (Average correctness)

$$Accuracy = \frac{1}{N} \sum_{i=1}^N \mathbb{1}_{\hat{y}_i = y_i}$$

## Classification: Confusion Matrix



## Decision Trees

- $2^{2^n}$  trees for  $n$  boolean attributes

### Entropy

$$I\left(\frac{p}{p+n}, \frac{n}{p+n}\right) = -\frac{p}{p+n} \log_2 \frac{p}{p+n} - \frac{n}{p+n} \log_2 \frac{n}{p+n}$$

### Remainder

$$remainder(A) = \sum_{i=1}^v \frac{p_i + n_i}{p + n} I\left(\frac{p_i}{p_i + n_i}, \frac{n_i}{p_i + n_i}\right)$$

### Information Gain

$$IG(A) = I\left(\frac{p}{p+n}, \frac{n}{p+n}\right) - remainder(A)$$

Entropy of this node      Entropy of children nodes

```
def DTL(examples, attributes, default):
    if examples is empty: return default
    if examples have the same classification:
        return classification
    if attributes is empty:
        return mode(examples)
    best = choose_attribute(attributes, examples)
    tree = a new decision tree with root best
    for each value v_i of best:
        examples_i = {rows in examples with best = v_i}
        subtree = DTL(examples_i, attributes - best, mode(examples))
        add a branch to tree with label v_i and subtree subtree
```

**Q:** How to deal with continuous-valued attributes?

**A:** Partition them into discrete sets of intervals

**Q:** What if some values are missing?

- A:**
- Assign the most common value of the attribute
  - Assign the most common value of the attribute with the same output
  - Assign probability to each possible value and sample
  - Drop the attribute
  - Drop the rows
  - ...

## Overfitting

- Occurs when algorithm fits training data too closely, resulting in worse performance on test data

## Occam's Razor

- Short/simple** hypotheses that fit data **unlikely to be coincidence**
- Long/complex** hypotheses that fit data may be **coincidence**

## Pruning

- Methods: Min samples leaf, max depth, ...
- Choose the majority at the node to be pruned

Core idea: Prevent overfitting by ignoring noise.